

Final report: Camera Based 2D Feature Tracking

Gia Minh Hoang

F.P1 Match 3D Objects

Lines 317 – 344 in `camFusion_Student.cpp` implements method “matchBoundingBoxes”. It iterates through all the keypoint matches as outer loop. Each keypoint match is checked whether it is enclosed in a bounding box for current and previous frames. Helper function `isEnclosedInBoundingBox` (lines 296 – 315) uses `cv::Rect.contains(cv::Point)` to find the bounding box. Moreover, if there is more than 1 bounding boxes, the helper function returns false to avoid mismatching.

A variable `match2DMatrix` whose each row represents the index of bounding box in current frame and each column shows that in previous frame counts the number of number of keypoint correspondences in order to find the bounding matches.

FP.2 Compute Lidar-based TTC

Lines 248 – 294 in `camFusion_Student.cpp` computes lidar-based TTC. As lidar cannot measure the velocity directly, it can only be done by differencing 2 successive measurements. For each measurement i.e., point cloud, the closest point in x-axis (forward direction) is found. It is compared with the mean value in x-axis of all the points scaled by a factor in order to make it robust to possible outliers. Then the TTC is calculated as

$$TTC_{lidar} = \frac{1}{frameRate} \frac{minXCurr}{minXPrev - minXCurr}, \quad (1)$$

where `minXCurr` and `minXPrev` are the closest points in x-axis in current and previous frame respectively. Note that the preceding vehicle is running at higher speed than the ego vehicle, `minXPrev < minXCurr` causes negative TTC. Thus, the function returns `TCC = NAN` and signals “No potential collision, preceding vehicle is running at higher speed!”.

FP.3 Associate Keypoint Correspondences with Bounding Boxes

Lines 133 – 194 in `camFusion_Student.cpp` adds the keypoint correspondings to the bounding box. It simply loops over the keypoint matches and uses `cv::Rect.contains(cv::Point)` similar to task FP.1. To make it robust to the outliers, I have tested 3 methods as follows:

1. Compute a mean of all the euclidean distances between keypoint matches and then remove those that are too far away from the mean.
2. Compute a median of all the euclidean distances between keypoint matches and then remove those that are too far away from the median.
3. Compute a mean and a standard deviation (1σ) of all the euclidean distances between keypoint matches and then remove those that are far away more than 3σ from the mean.

I found that method 3 did not work well as some outliers made the standard deviation so large. Methods 1 and 2 worked better. Though median is slightly better than mean, I found median method

did not work for the case HARRIS/BRISK due to the fact that HARRIS detects only the few points and the median are rejected all the points.

FP.4 Compute Camera-based TTC

Lines 197 – 245 in `camFusion_Student.cpp` computes camera-based TTC. As monocular cameras are not able to measure metric distances, TTC is computed by observing relative height change (also called scale change) directly in the image. Identifiable keypoints are detected and tracked from one frame to the next, the distance between all keypoints on the vehicle relative to each other to compute a robust estimate of the height ratio in our TTC equation as

$$TTC_{camera} = \frac{-1}{frameRate} \frac{1}{1 - medDistRatio} \quad (2)$$

where *medDistRatio* is the median of all the relative distance ratios computed from keypoint matches in 2 successive images to remove the outliers. The *medDistRatio* is calculated by iterating through all points in the `vector<cv::DMatch> kptMatches`, and each of these points to all other points in the same vector using an inner for loop.

FP.5 Performance Evaluation 1

The lidar-based TTC is logged using all possible combinations (of detectors and descriptors). Since AKAZE descriptor only works with AKAZE detector and descriptor ORB cannot function with detector SIFT, we have 35 combinations.

The results are collected in the `results.csv` in directory `/SFND_3D_Feature_Tracking/report/`

The keypoint matches provided by all combinations (of detectors and descriptors) are able to find the bounding box matches. Thus, the lidar-based TTC estimates are the same regardless of the combinations. Let take a look at an example with SHITOMASI/BRISK as follows:

Table 1: Results syntheses for SHITOMASI/BRISK.

Detector	Descriptor	Image id	No matched keypoints in ROI	TTC lidar	TTC camera
SHITOMASI	BRISK	0	0	nan	nan
SHITOMASI	BRISK	1	59	12.9722	14.1119
SHITOMASI	BRISK	2	59	12.264	12.9876
SHITOMASI	BRISK	3	58	13.9161	13.4904
SHITOMASI	BRISK	4	63	7.11572	12.398
SHITOMASI	BRISK	5	63	16.2511	12.6754
SHITOMASI	BRISK	6	61	12.4213	14.5029
SHITOMASI	BRISK	7	64	34.3404	12.9117
SHITOMASI	BRISK	8	58	9.34376	15.5997
SHITOMASI	BRISK	9	66	18.1318	11.5126
SHITOMASI	BRISK	10	57	18.0318	14.685
SHITOMASI	BRISK	11	60	3.83244	11.3024
SHITOMASI	BRISK	12	78	nan	11.6524
SHITOMASI	BRISK	13	75	9.22307	11.7243
SHITOMASI	BRISK	14	75	10.9678	11.5069

SHITOMASI	BRISK	15	64	8.09422	9.33785
SHITOMASI	BRISK	16	63	3.17535	11.3079
SHITOMASI	BRISK	17	66	nan	11.4487
SHITOMASI	BRISK	18	71	8.30978	9.12792

TTC = NAN for the first image is normal because 2 lidar measurements are needed to compute TTC.

Let consider the second NAN for image 12, the reason is that the preceding vehicle has higher speed than the ego vehicle leading to increased relative distance between them as explained in task FP.2. The top view perspective of the Lidar points in the following figure confirms this observation.




Figure 1: Top view perspective of the Lidar points for frame 11 (top) and frame 12 (bottom). Note that the relative distance increases.

Similarly, TTC = NAN for image 17 has the same reason as illustrated by the figure below.





Figure 2: Top view perspective of the Lidar points for frame 16 (top) and frame 17 (bottom). Note that the relative distance increases.

Finally, let try to understand why the TCC jumped to rather high value of 34.34 s in image 7. The front views with displayed bounding box, lidar points, and keypoints as well as top view perspective of the lidar points are depicted in the following figure.






Figure 3: Front views displayed bounding box, lidar points, and keypoints for image 6 (top left) and 7 (bottom left) while top view perspective of the lidar points for image 6 (top right) and 7 (bottom right).

Based on the top view of lidar point, using equation (1) lidar-based TTC = 25.17 s leading to an error of 36%. The possible reason could be some lidar points are unstable to use in TTC computation especially low speed scenario when differing position to get velocity is prone to error.

FP.6 Performance Evaluation 2

The results is stored in the results.csv in directory /SFND_3D_Feature_Tracking/report/

Based on the results.csv, there are a lot of combinations (of detectors and descriptors) that yields rather stable lidar-based TTC.

However, in midterm project, the top 3 detector/descriptor combinations are FAST/BRIEF, FAST/BRISK, and FAST/ORB which are the fastest in execution time while maintaining rather good portion of keypoints on the preceding vehicle, and matching rather good portion of keypoints between successive images. Let consider their camera-based TTC estimates in the following table. Overall, TTC estimates are rather stable except FAST/BRISK has 1 irrelevant estimate.

Table 2: Results syntheses for FAST/BRIEF, FAST/BRISK, and FAST/ORB.

Detector	Descriptor	Image id	No matched keypoints in ROI	TTC lidar	TTC camera
FAST	BRISK	0	0	nan	nan
FAST	BRISK	1	58	12.9722	12.3
FAST	BRISK	2	67	12.264	12.3453
FAST	BRISK	3	75	13.9161	14.1877
FAST	BRISK	4	68	7.11572	12.8857
FAST	BRISK	5	65	16.2511	-Inf
FAST	BRISK	6	69	12.4213	13.0386
FAST	BRISK	7	77	34.3404	12.041
FAST	BRISK	8	72	9.34376	11.4066
FAST	BRISK	9	75	18.1318	11.8684
FAST	BRISK	10	74	18.0318	13.3473
FAST	BRISK	11	67	3.83244	12.9492
FAST	BRISK	12	86	nan	12.1174
FAST	BRISK	13	87	9.22307	12.1074
FAST	BRISK	14	82	10.9678	11.6077
FAST	BRISK	15	79	8.09422	11.4079
FAST	BRISK	16	86	3.17535	12.2566
FAST	BRISK	17	89	nan	9.2933
FAST	BRISK	18	84	8.30978	11.8606
FAST	BRIEF	0	0	nan	nan

FAST	BRIEF	1	83	12.9722	11.1897
FAST	BRIEF	2	85	12.264	12.5319
FAST	BRIEF	3	92	13.9161	17.3793
FAST	BRIEF	4	88	7.11572	13.6612
FAST	BRIEF	5	84	16.2511	30.2446
FAST	BRIEF	6	81	12.4213	13.5856
FAST	BRIEF	7	97	34.3404	12.2781
FAST	BRIEF	8	82	9.34376	12.5357
FAST	BRIEF	9	87	18.1318	13.4042
FAST	BRIEF	10	88	18.0318	13.4879
FAST	BRIEF	11	83	3.83244	14.132
FAST	BRIEF	12	93	nan	12.6994
FAST	BRIEF	13	94	9.22307	12.3705
FAST	BRIEF	14	104	10.9678	11.7034
FAST	BRIEF	15	97	8.09422	12.1403
FAST	BRIEF	16	102	3.17535	12.5062
FAST	BRIEF	17	104	nan	8.48654
FAST	BRIEF	18	96	8.30978	12.6126
<hr/>					
FAST	ORB	0	0	nan	nan
FAST	ORB	1	75	12.9722	12.0985
FAST	ORB	2	84	12.264	12.8193
FAST	ORB	3	89	13.9161	16.3163
FAST	ORB	4	84	7.11572	13.9934
FAST	ORB	5	82	16.2511	37.8535
FAST	ORB	6	78	12.4213	14.0835
FAST	ORB	7	86	34.3404	12.559
FAST	ORB	8	90	9.34376	11.9141
FAST	ORB	9	93	18.1318	12.8711
FAST	ORB	10	89	18.0318	13.8572
FAST	ORB	11	78	3.83244	18.3152
FAST	ORB	12	87	nan	12.7731
FAST	ORB	13	94	9.22307	12.9215
FAST	ORB	14	96	10.9678	11.1115
FAST	ORB	15	97	8.09422	10.8439
FAST	ORB	16	100	3.17535	11.3939
FAST	ORB	17	102	nan	10.6085
FAST	ORB	18	95	8.30978	12.6746

The camera-based TTC estimates are worse if using HARRIS detector as depicted in the following table as an example of HARRIS/BRISK. The reason is that HARRIS detector provides less keypoints than other detectors making the TTC computation is not reliable.

Table 3: Results syntheses for HARRIS/BRISK.

Detector	Descriptor	Image id	No matched keypoints in ROI	TTC lidar	TTC camera
HARRIS	BRISK	0	0	nan	nan
HARRIS	BRISK	1	7	12.9722	10.9082
HARRIS	BRISK	2	8	12.264	10.586
HARRIS	BRISK	3	7	13.9161	-80.8525
HARRIS	BRISK	4	11	7.11572	11.5792
HARRIS	BRISK	5	12	16.2511	-Inf
HARRIS	BRISK	6	9	12.4213	12.9945
HARRIS	BRISK	7	15	34.3404	11.6947
HARRIS	BRISK	8	14	9.34376	17.6204
HARRIS	BRISK	9	7	18.1318	nan
HARRIS	BRISK	10	3	18.0318	-Inf
HARRIS	BRISK	11	6	3.83244	-Inf
HARRIS	BRISK	12	8	nan	12.245
HARRIS	BRISK	13	12	9.22307	568.322
HARRIS	BRISK	14	9	10.9678	7.72144
HARRIS	BRISK	15	9	8.09422	-13.6263
HARRIS	BRISK	16	17	3.17535	6.65726
HARRIS	BRISK	17	7	nan	12.5848
HARRIS	BRISK	18	4	8.30978	-Inf

For example, the following figure shows that HARRIS produced only 12 matched keypoints but some of them are located on the farther vehicle on the left lane.




Figure 4: Front views displayed bounding box, lidar points, and keypoints using HARRIS/BRISK for image 5.